

Real-time Graphics

1. Graphics Pipeline, Shaders

Martin Samuelčík

Juraj Starinský

Graphics Pipeline

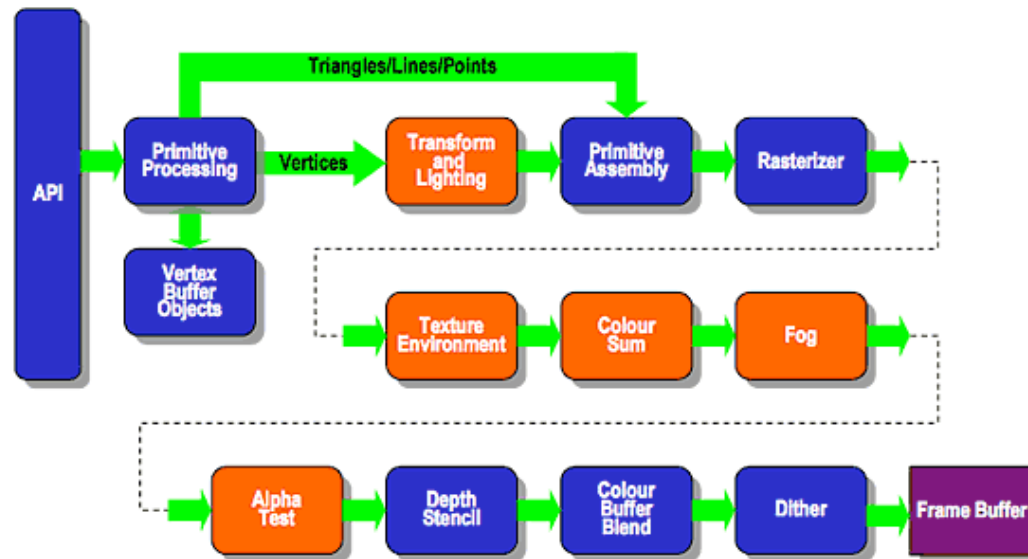
- Rasterization-based real-time rendering
- Supported by common hardware - graphics cards
- Input = 3D representation of scene
- Output = 2D raster image
- Stream processing
- Another ways of rendering: raytracing, global illumination (radiosity), REYES, ...



OpenGL Graphics Pipeline



- Fixed pipeline - prerequisite for this course
- Knowledge of extensions mechanism
- Some parts are depreciated

Existing Fixed Function Pipeline




First generation

- Vertex: transform, clip, project
- Pixel: color interpolation of lines
- Frame buffer: overwrite
- Dates: prior to 1987



Second generation

- Vertex: lighting generation
- Pixel: depth interpolation, triangles
- Frame buffer: depth buffer, blending
- Dates: 1987-1992



Real-time Graphics

Martin Samuelčík, Juraj Starinský

Third generation

- Vertex: texture coordinate transformation
- Pixel: texture coordinate interpolation, texture evaluation and filtering
- Dates: 1992 - 2000



Real-time Graphics

Martin Samuelčík, Juraj Starinský

Fourth generation

- Programmable shading: Vertex, Pixel, Geometry
- Multi GPU (SLI, Crossfire), Full floating point, GPGPU



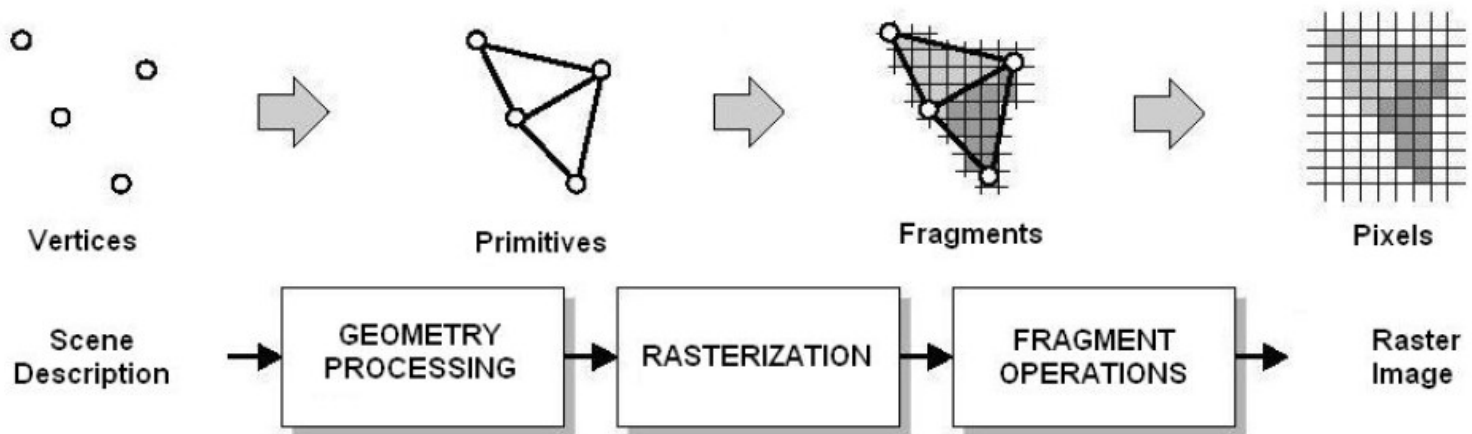
Fifth generation

- Multi cores
- Merging CPU and GPU
 - IBM Cell (8+1cores)
 - AMD Fusion (CPU+GPU)
 - Intel Larabee – canceled
- Programming
 - Tessellation shader
 - DX11 (Compute shader)
 - CUDA, AMD Brook
 - OpenCL



Parts of pipeline

- Application – user settings, scene description
- Geometry – transformations, clipping, projection
- Rasterization – computations, texturing
- Fragments – tests, blending, operations




Application stage

- Scene description
- User full control
- Scene complexity – LOD, clipping
- Scene dynamics (physics)
- Handling – mouse, keyboard



Geometry stage

- Per-primitive operations
 - Model & view transform
 - Per-vertex lightning & shading
 - Projection
 - Clipping
 - Screen mapping



Fragment tests

- Ownership
- Scissor test
- Alpha test
- Stencil test
- Depth test
- Alpha blending
- Logical operations



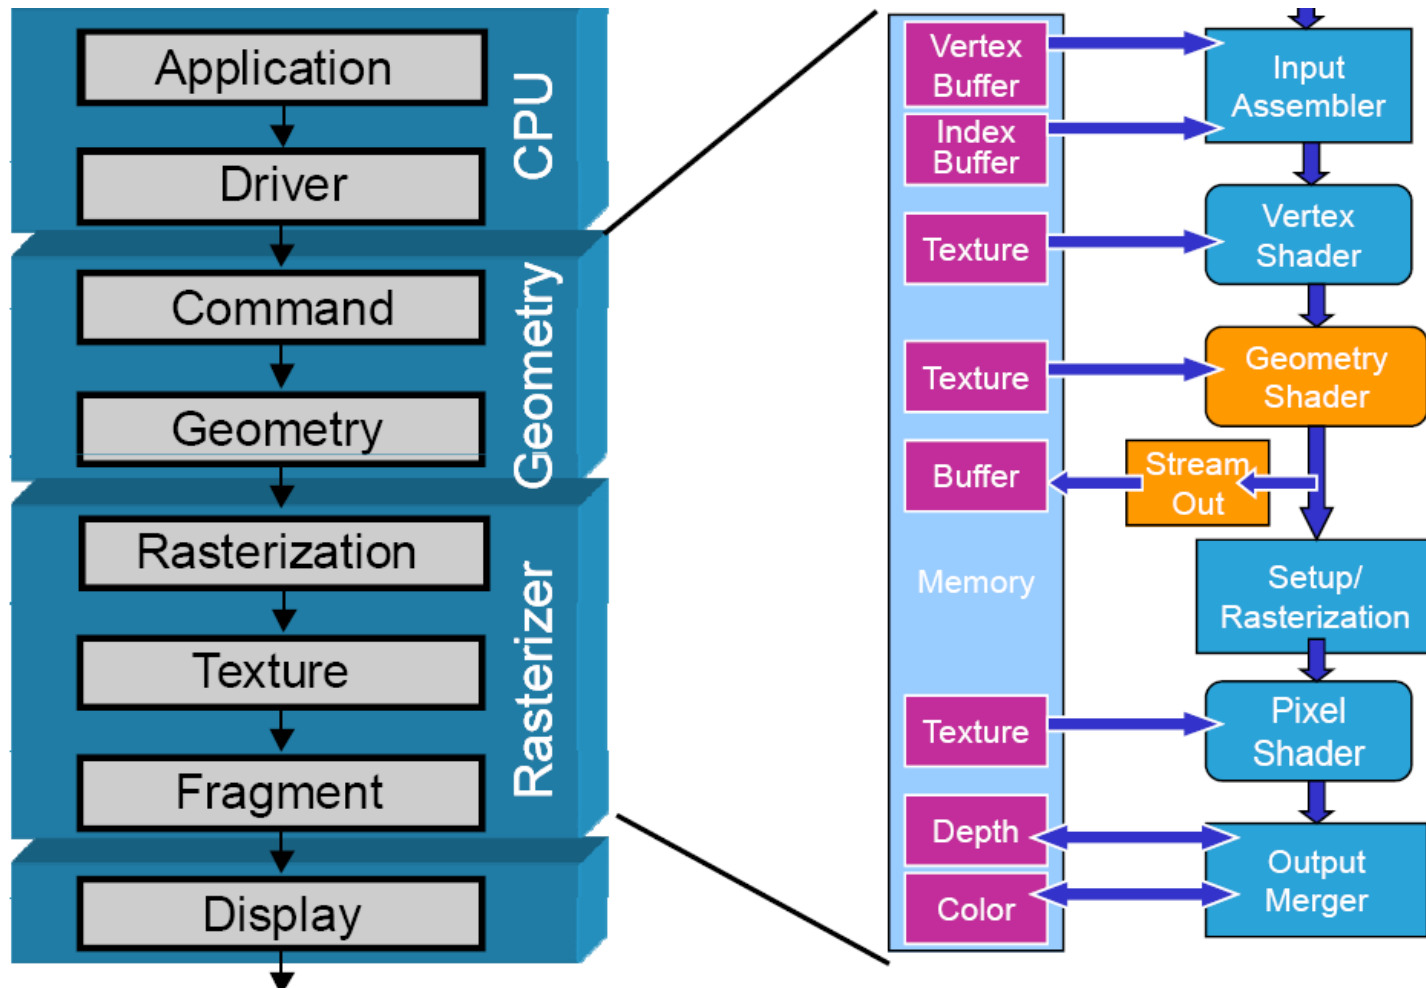
A simple three dimensional scene



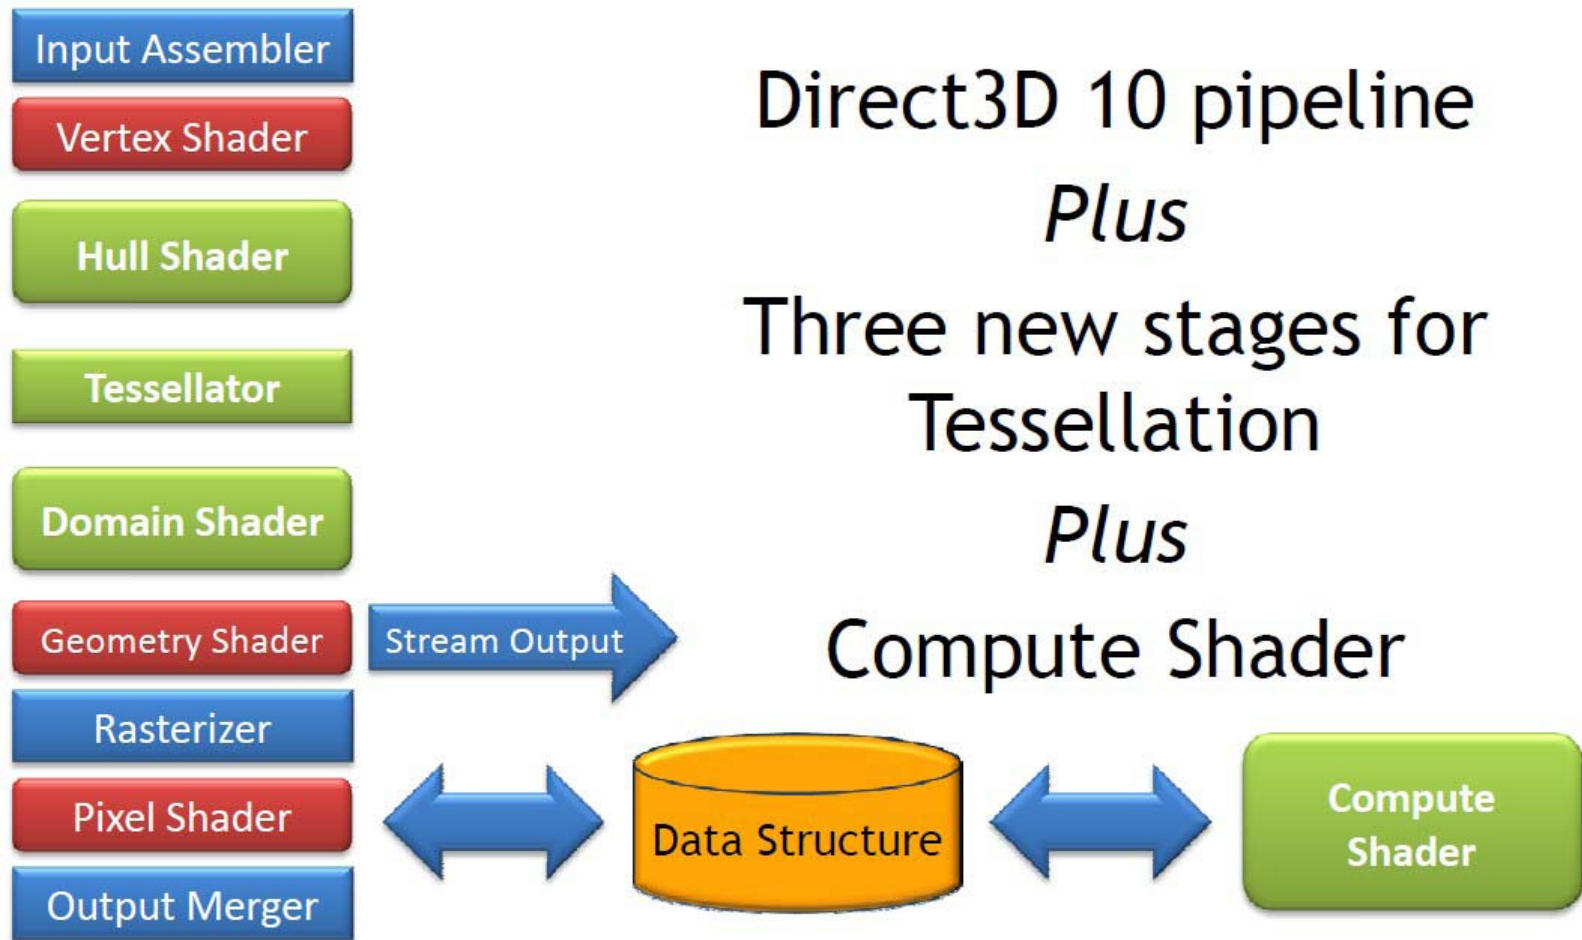
Z-buffer representation



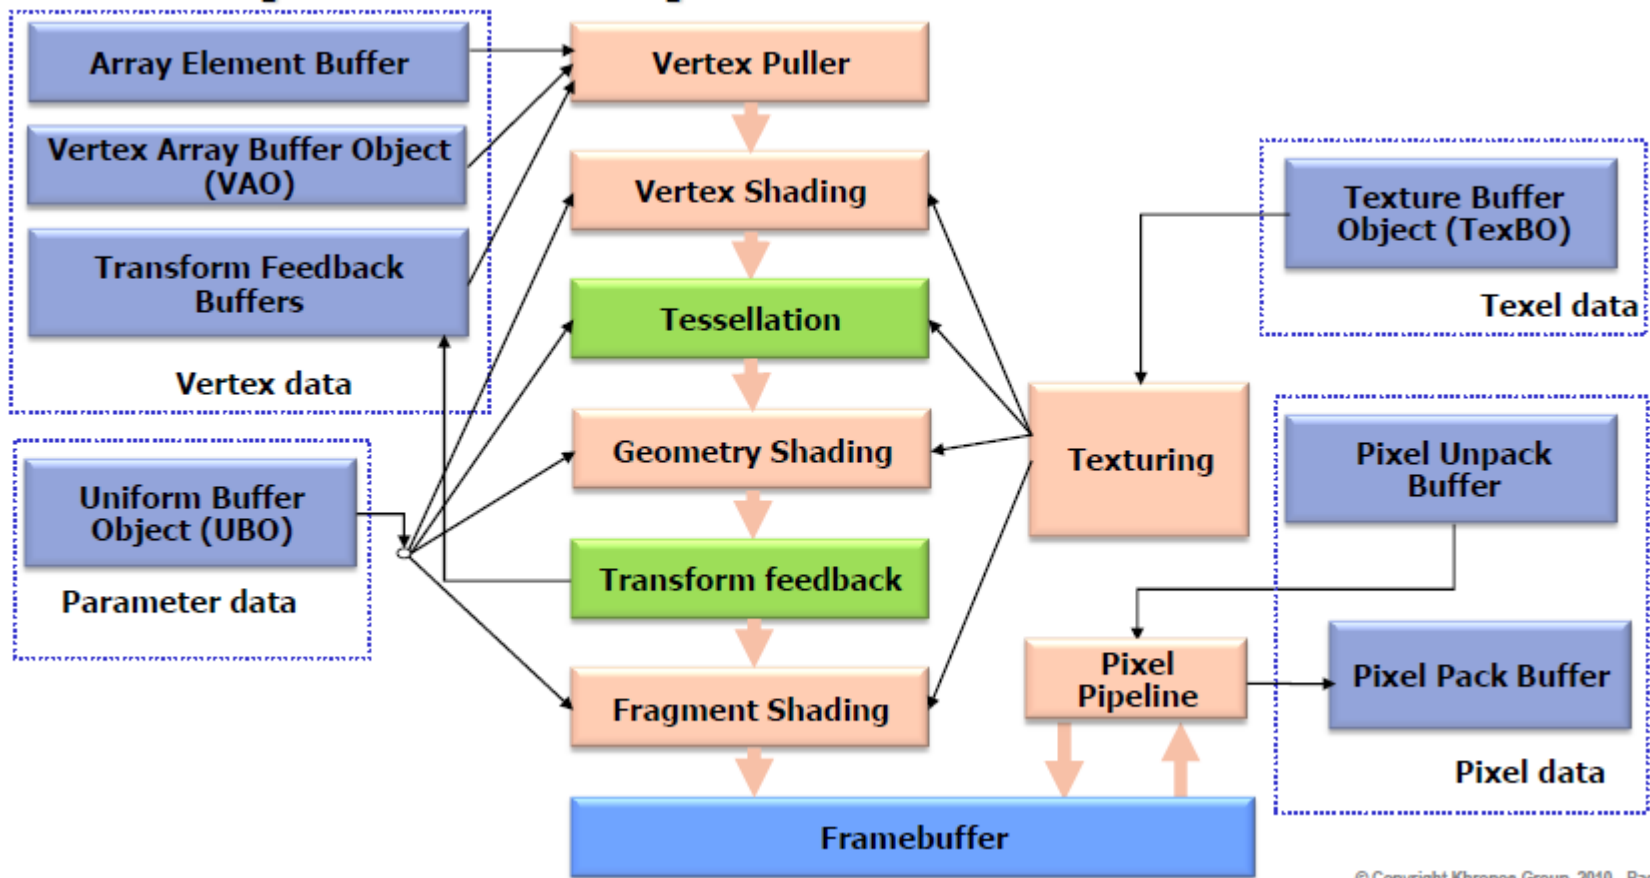
DirectX 10 pipeline



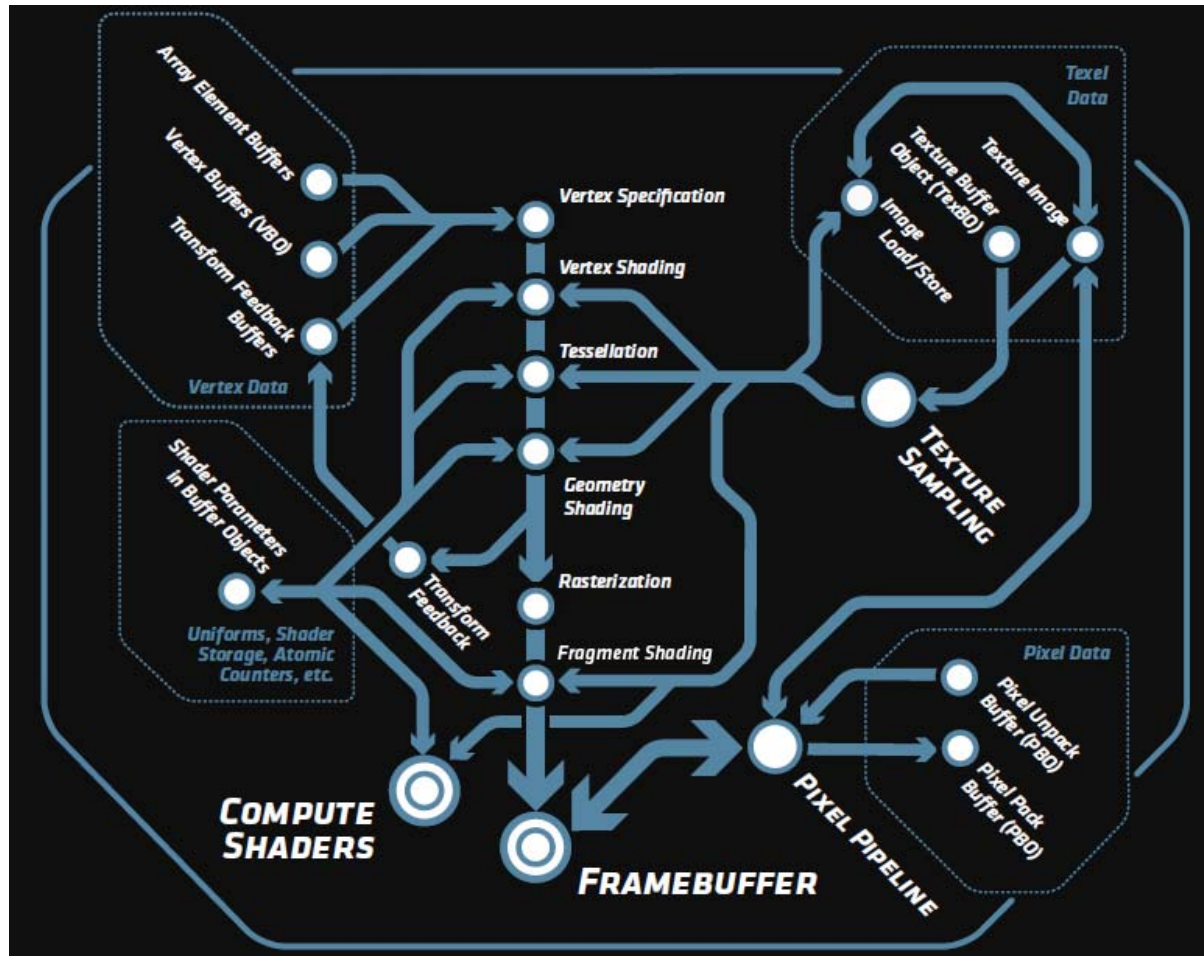
DirectX 11 pipeline



OpenGL 4.0 pipeline




OpenGL 4.3 pipeline



Shaders

- Programs for creating new geometry, changing vertices and shaders
- Take control of processing on GPU
- Move some computation from CPU to GPU




Shader languages

- Assembler - different for ATI, Nvidia
- Source code shaders - ATI
- Cg – C for graphics (Nvidia), HLSL (DirectX)
- GLSL - OpenGL Shading Language
 - OpenGL 2.0 standard
 - ATI, Nvidia



Cg

- Nvidia extensions
- Various profiles
- Compiler, runtime libraries
- Examples – Cg Toolkit
- http://developer.nvidia.com/object/cg_toolkit.html



Cg shaders

```
struct input_data
{
    float4 position    : POSITION;
};

struct output_data
{
    float4 position    : POSITION;
};

output_data main(input_data IN)
{
    output_data OUT;
    OUT.position = IN.position;
    return OUT;
}
```

Vertex shader

```
struct input_data{
    float2 tc : TEXCOORD0;
};
struct output_data{
    float4 color : COLOR;
};


output_data main(input_data IN,
                  uniform sampler2D textureIn )
{
    output_data OUT;
    OUT.color = tex2D(textureIn, IN.tc);
    return OUT;
}
```

Fragment shader



GL shading language

- ANSI C-like language for writing shaders
- <http://www.opengl.org/registry/doc/GLSLangSpec.4.10.6.clean.pdf>
- Extended with mechanisms from C++ and vector and matrix types
- Part of core specification
- Older functionality is depreciated



OpenGL extension for GLSL

- New language for writing shaders
 - `GL_ARB_shading_language_100`, ...
- New shader programs
 - `GL_ARB_fragment_shader`
 - `GL_ARB_vertex_shader`
 - `GL_ARB_geometry_shader4`
 - `GL_ARB_tessellation_shader`, ...
- Management using shader objects
 - `GL_ARB_shader_objects`
- <http://www.opengl.org/registry/>



GLSL basic types

- void
- float vec2 vec3 vec4
- mat2 mat3 mat4
- int ivec2 ivec3 ivec4
- bool bvec2 bvec3 bvec4
- sampler2D samplerCube
- samplerShadow2D



GLSL type qualifiers

- Const – compile-time constant
- Attribute – for passing data for vertex
- Uniform – global variable, read-only
- Varying – vertex->fragment sh. data
- In – parameters passed into function
- Out – passed out of function
- Inout – in & out of function



GLSL scalar constructors

- `int(bool)` // converts a Boolean value to an int
- `int(float)` // converts a float value to an int
- `float(bool)` // converts a Boolean value to a float
- `float(int)` // converts an integer value to a float
- `bool(float)` // converts a float value to a Boolean
- `bool(int)` // converts an integer value to a Boolean



GLSL matrix constructors

```
vec3(float)      // initializes each component of a vec3 with the float
vec4(ivec4)      // makes a vec4 from an ivec4, with component-wise conversion

vec2(float, float)      // initializes a vec2 with 2 floats
ivec3(int, int, int)    // initializes an ivec3 with 3 ints
bvec4(int, int, float, float) // initializes with 4 Boolean conversions

vec2(vec3) // drops the third component of a vec3
vec3(vec4) // drops the fourth component of a vec4

vec3(vec2, float) // vec3.x = vec2.x, vec3.y = vec2.y, vec3.z = float
vec3(float, vec2) // vec3.x = float, vec3.y = vec2.x, vec3.z = vec2.y
vec4(vec3, float)
vec4(float, vec3)
vec4(vec2, vec2)

mat2(float)
mat3(float)
mat4(float)

mat2(vec2, vec2);
mat3(vec3, vec3, vec3);
mat4(vec4, vec4, vec4, vec4);

mat2(float, float,
      float, float);

mat3(float, float, float,
      float, float, float,
      float, float, float);

mat4(float, float, float, float,
      float, float, float, float,
      float, float, float, float,
      float, float, float, float);
```



GLSL vectors & matrices

- Component access:
 - $\{x,y,z,w\}$, $\{r,g,b,a\}$, $\{s,t,p,q\}$
 - Allows access to multiple components
 - Allows access using indexing []

```
vec4 pos = vec4(1.0, 2.0, 3.0, 4.0);  
pos.xw = vec2(5.0, 6.0);      // pos = (5.0, 2.0, 3.0, 6.0)  
pos.wx = vec2(7.0, 8.0);      // pos = (8.0, 2.0, 3.0, 7.0)  
pos.xx = vec2(3.0, 4.0);      // illegal - 'x' used twice  
pos.xy = vec3(1.0, 2.0, 3.0); // illegal - mismatch between vec2 and vec3
```

- Simple addition and multiplication:
 - $\text{vec3 } u, v, w; w = u + v; w = u * v;$
 - $\text{mat3 } m, n, o; m = n * o; m = n + o$



GLSL built-in functions

- Angle and Trigonometry (radians, degrees, sin, cos, tan, asin, acos, ...)
- Exponential (pow, exp, log, sqrt, ...)
- Math (abs, floor, mod, min, max, clamp, mix, step, mod, ...)
- Geometric (length, distance, dot, cross, normalize, ftransform, reflect, ...)
- Matrix (matrixCompMult)
- Vector Relational (lessThan, greaterThan, equal, any, all, ...)
- Texture Lookup (textureD, textureDLod, textureDProj, textureCube, shadowD, ...)
- Noise (noise1, noise2, ...)
- Fragment processing (discard, dFdx, dFdy, fwidth, ...)




GLSL built-in variables

- For access of data
- Part from OpenGL state
- User defined data
- Variables for input and output of shaders
- Based on fixed functionality pipeline



GLSL VS built-in variables



GLSL GS built-in variables

- Input:

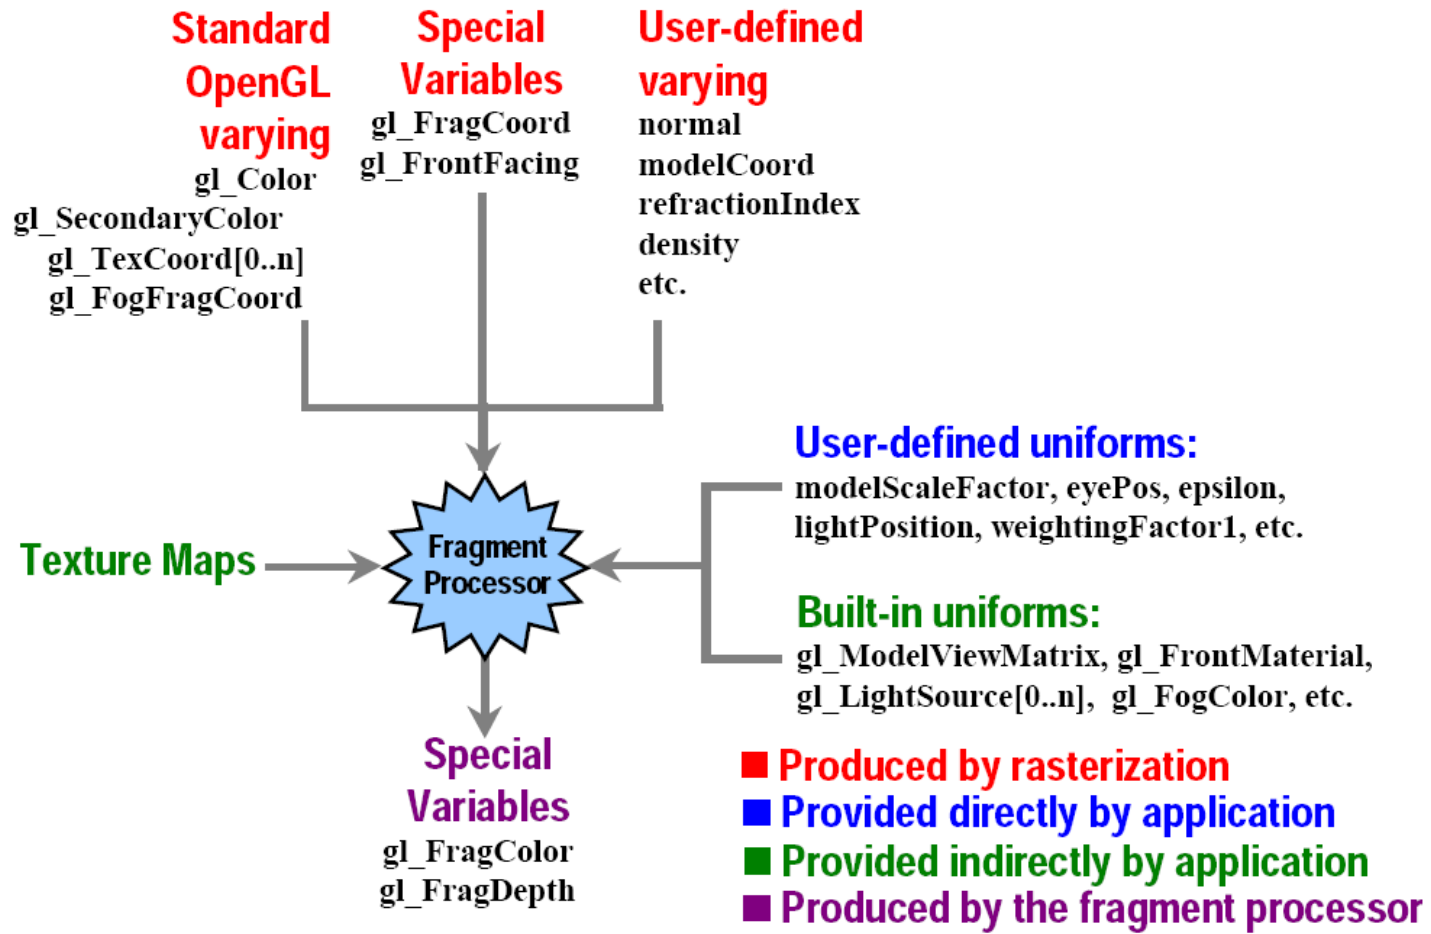
- varying in vec4 gl_FrontColorIn[gl_VerticesIn]
- varying in vec4 gl_BackColorIn[gl_VerticesIn]
- varying in vec4 gl_FrontSecondaryColorIn[gl_VerticesIn]
- varying in vec4 gl_BackSecondaryColorIn[gl_VerticesIn]
- varying in vec4 gl_TexCoordIn[gl_VerticesIn][]
- varying in vec4 gl_PositionIn[gl_VerticesIn]
- varying in float gl_PointSizeIn[gl_VerticesIn]
- varying in vec4 gl_ClipVertexIn[gl_VerticesIn]

- Output:

- varying vec4 gl_FrontColor
- varying vec4 gl_BackColor
- varying vec4 gl_FrontSecondaryColor
- varying vec4 gl_BackSecondaryColor
- varying vec4 gl_TexCoord[]
- varying out vec4 gl_FrontColor
- varying out vec4 gl_BackColor
- varying out vec4 gl_FrontSecondaryColor
- varying out vec4 gl_BackSecondaryColor
- varying out vec4 gl_TexCoord[];



GLSL FS built-in variables



GLSL – vertex, fragment

```
const vec4 AMBIENT = vec4( 0.9, 0.9, 0.1, 1.0 );
const vec4 SPECULAR = vec4( 1.0, 1.0, 1.0, 1.0 );
uniform vec4 light;

varying vec4 Ca;
varying vec4 Cd;
varying vec4 Cs;

varying vec4 V_eye;
varying vec4 L_eye;
varying vec4 N_eye;

void main(void)
{
    V_eye = gl_ModelViewMatrix * gl_Vertex;
    L_eye = (gl_ModelViewMatrix * light) - V_eye;
    N_eye = vec4(gl_NormalMatrix * gl_Normal, 1.0);

    gl_Position = gl_ProjectionMatrix * V_eye;
    V_eye = -V_eye;

    Ca = AMBIENT;
    Cd = gl_Color;
    Cs = SPECULAR;
}
```

```
varying vec4 Ca;
varying vec4 Cd;
varying vec4 Cs;

varying vec4 V_eye;
varying vec4 L_eye;
varying vec4 N_eye;

vec3 reflect(vec3 N, vec3 L)
{
    return 2.0*N*dot(N, L) - L;
}

void main(void)
{
    vec3 V = normalize(vec3(V_eye));
    vec3 L = normalize(vec3(L_eye));
    vec3 N = normalize(vec3(N_eye));

    float diffuse = clamp(dot(L, N), 0.0, 1.0);

    vec3 R = reflect(N, L);
    float specular = pow(clamp(dot(R, V), 0.0, 1.0), 16);

    gl_FragColor = Ca + (Cd*diffuse) + (Cs*specular);
}
```



Geometry shader example

```
#version 120
#extension GL_EXT_geometry_shader4: enable
uniform float FpNum;
void main()
{
    int num = int( FpNum + 0.99 );
    float dt = 1. / float(num);
    float t = 0.;
    for( int i = 0; i <= num; i++ )
    {
        float omt = 1. - t;
        float omt2 = omt * omt;
        float omt3 = omt * omt2;
        float t2 = t * t;
        float t3 = t * t2;
        vec4 xyzw = omt3 * gl_PositionIn[0].xyzw +
        3. * t * omt2 * gl_PositionIn[1].xyzw +
        3. * t2 * omt * gl_PositionIn[2].xyzw +
        t3 * gl_PositionIn[3].xyzw;
        gl_Position = xyzw;
        EmitVertex()
        t += dt;
    }
    EndPrimitive();
}
```

```
void main()
{
    gl_Position = gl_ModelViewProjectionMatrix * gl_Vertex;
}

void main()
{
    gl_FragColor = vec4( 0., 1., 0., 1. );
}
```




Tessellation shaders

- Works on patch – given by set of vertices and per-patch attributes
- Tessellation control shader transforms per-vertex data and per-patch attr.
- Tessellator decomposes patch into set of new primitives based on tess. level
- Tessellation evaluation shader computes position and attributes of new generated vertices



Tessellation shaders



Shaders management

- Shader objects – shaders with unique identifier
- Creation: *glCreateShaderObject(type)*, type:
 - GL_VERTEX_SHADER
 - GL_GEOMETRY_SHADER
 - GL_FRAGMENT_SHADER, ...
- Setting source: *glShaderSource(shaderID, numStrings, strings, length)*
- Compilation: *glCompileShader(shaderID)*



Shaders management

- Shader programs - container for shader objects, set of shaders that are linked together
- Creation: `prog = glCreateProgramObject()`
- Adding shader: `glAttachObject(programID, shaderID)`
- Linking: `glLinkProgram(programID)`
- Set as current:
`glUseProgramObject(programID)`



Management example

```
GLhandle g_programObj;  
GLhandle g_vertexShader;  
GLhandle g_fragmentShader;  
  
g_vertexShader = glCreateShaderObjectARB( GL_VERTEX_SHADER );  
unsigned char *vertexShaderAssembly = readShaderFile( "vertex_shader.vert" );  
vertexShaderStrings[0] = (char*)vertexShaderAssembly;  
glShaderSource( g_vertexShader, 1, vertexShaderStrings, NULL );  
glCompileShader( g_vertexShader );  
delete vertexShaderAssembly;  
  
g_fragmentShader = glCreateShaderObject( GL_FRAGMENT_SHADER );  
unsigned char *fragmentShaderAssembly = readShaderFile( "fragment_shader.frag" );  
fragmentShaderStrings[0] = (char*)fragmentShaderAssembly;  
glShaderSource( g_fragmentShader, 1, fragmentShaderStrings, NULL );  
glCompileShader( g_fragmentShader );  
delete fragmentShaderAssembly;  
  
g_programObj = glCreateProgramObject();  
glAttachObject( g_programObj, g_vertexShader );  
glAttachObject( g_programObj, g_fragmentShader );  
  
glLinkProgram( g_programObj );  
glGetObjectParameteriv( g_programObj, GL_OBJECT_LINK_STATUS, &bLinked );
```



Passing variables


- From application to shaders, based on location of variable in shader program:
 - Glint `glGetAttribLocation`(GLhandle program, const GLchar* name);
 - Glint `glGetUniformLocation`(GLhandle program, const GLchar * name);
 - void `glUniform`{1|2|3|4}{f|i}(GLint location, TYPE val);
 - void `glUniform`{1|2|3|4}{f|i}v(GLint location, GLuint count, const TYPE * vals);
 - void `glUniformMatrix`{2|3|4}{f}v(GLint location, GLuint count, GLboolean transpose, const GLfloat * vals);
 - void `glVertexAttrib`{1|2|3|4}{s|f|d}(GLuint index, TYPE val);
 - void `glVertexAttrib`{1|2|3|4}{s|f|d}v(GLuint index, const TYPE * vals);
- Possibility to sent array of attributes or uniforms



Unified architecture

- Consistent instruction set across all shader types
- Flexible use of the graphics hardware

Why unify?





Heavy Geometry
Workload Perf = 4


Why unify?



Heavy Geometry
Workload Perf = 12



Heavy Pixel
Workload Perf = 8



Heavy Pixel
Workload Perf = 12



Real-time Graphics

Martin Samuelčík, Juraj Starinský

Tools

- Debugging GLSL: your way, gDEBugger, glslDevil, Fx Composer(+ Shader Debugger)
- Extensions: GLEE, GLEW
- Render Monkey -
<http://developer.amd.com/archive/gpu/rendermonkey/pages/default.aspx>
- Shader Designer -
<http://www.opengl.org/sdk/tools/ShaderDesigner/>
- Books:
<http://www.opengl.org/documentation/books/>



Questions?

